



COMP 3200

Artificial Intelligence

Lecture 3

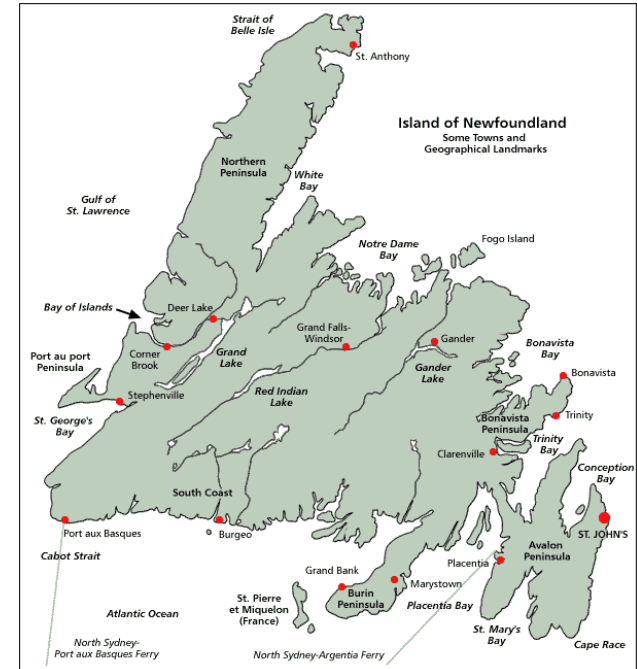
Problem Solving and Search

Problem Solving Agents

- Rational agents maximize performance
- Simplest Case: Satisfy a goal condition
 - Reach the end of a maze
 - Solve a puzzle
- Problem Solving Agent
 - Type of Goal-Based Agent
 - Find a sequence of actions that satisfies goal

Example Scenario

- Tourist comes to NL
 - Lands in St. John's
- Wants to go to one of:
 - St. Anthony
 - Port aux Basques
- Extra constraints?
 - Has X hours total
 - Has \$\$ total for gas




Goal Formulation

- Goal is a **set of states** in the environment
 - [St. Anthony, Port aux Basques]
- Goal is **satisfied** if one of them is reached
 - If a sequence of actions leads us to a goal state, the goal has been **satisfied**, and the **problem is solved**
- **Goal** can also be a function evaluation
 - If (enemy health ≤ 0)
 - If (enemy in checkmate)
 - If (player has all pieces of triform AND Ganon dead)

Problem Formulation

- What states to consider?
 - Every square meter in Newfoundland
 - Towns in Newfoundland
- What actions to consider?
 - Move left by a meter
 - Move from town X to town Y
- What information is stored in a node?
 - Current town location
 - Distance / time driven so far



Well-Defined Problem

1. Initial State that the agent starts in
2. Actions possible from each State
 - Transition / Successor Function
3. The set of goal states
 - Goal test function
4. Action cost function
 - Path cost = $\text{sum}(\text{cost}(\text{actions in path}))$

Example Problem Definition

1. Initial State
 - St. John's
2. Actions
 - Towns connected by roads
3. Goal
 - [SA, PAB]
4. Action Cost Function
 - Travel distance or time




Problem Solution

- **Path** from start state to a goal state
 - Path = ordered sequence of actions
- **Solution cost** = path cost
 - Sum cost of actions (usually)
 - Most problems assume costs > 0
- **Optimal solution** has lowest path cost among all possible solutions


Example Graph Problem

- Initial State
 - State A
- Actions
 - (A,B) legal if edge
- Goal State
 - State C
- Action Cost Function
 - Edge label (A,B)=6




Example Graph Problem

- Fully Observable
- Static
- Discrete
- Deterministic
- Single Agent
- Sequential



Example Graph Problem

- Objective
 - Path to goal possible?
 - Shortest path to goal?
- Algorithm
 - Which to use?
- Let's try **search**





What is Search?

- Agent has several available actions
- Agent can explore various sequences of those actions
- Agent chooses the best sequence found
- This process in general is called **search**

Search

- Search will explore the problem state space
- Searching this space generates a **search tree**
- The search tree has **nodes** and **edges**


Search Tree





Search Tree

Expand



A





Search Tree





Search Tree





Search Tree





Search Tree





Search Tree



Search Tree





Search Tree





Continue Until
Goal Found

Search Tree





State Space: Finite

Search Tree




State Space: Finite
Search Tree: Infinite



Example: 15 Puzzle



Example: 15 Puzzle



Search Strategy





Search Tree Node Data

- **State**: The state in the graph
- **Parent Node**: Tree node that generated this node
- **Action**: Action that produced this node
- **Path Cost**: Cost of path so far from start
 - Traditionally denoted by $g(n)$
- **Depth**: Number of actions to this node
- Other: Can contain precomputed data (optimization)
 - Legal actions, heuristics, etc

Node vs. State

- Important distinction
- State
 - Configuration of the environment
- Node
 - Bookkeeping data structure
 - Exists only within the search tree
 - Nodes are 'on paths' in the search tree

The Fringe



General Uninformed Tree Search

1. Function **Tree-Search**(problem, strategy)
2. fringe = {Node(problem.initial_state)}
3. **while** (true)
4. if (fringe.empty) **return** fail
5. node = strategy.select_node(fringe)
6. if (node.state is goal) **return** solution
7. else fringe.add(Expand(node, problem))

Node Expansion

1. Function **Expand**(node, problem)
2. successors = { }
3. for a : problem.actions(node)
4. s = Node(node.state.do_action(a))
5. s.parent = node
6. s.g = node.g + s.action.cost
7. s.action = a
8. s.depth = node.depth + 1
9. successors.add(s)
10. **return** successors

Problem Solving Performance

- **Completeness**
 - Is it guaranteed to find a solution if it exists?
- **Optimality**
 - Does it find the optimal solution?
- **Time Complexity**
 - How long does it take to find a solution?
- **Space Complexity**
 - How much memory is needed to run the search?

Time + Space Complexity


- Measured on **size of input problem** / graph
- Time measures **number of nodes generated**
- Space measured by how much memory is needed to store the **maximum number of nodes** (worst case) needed at any given time
- Search / AI traditionally measures:
 - Branching factor **b** (max successors at any node)
 - Depth **d** (of the shallowest goal node)
 - Tree size $\sim b^d$

Heuristic Search

- Improve upon uninformed tree search
- Incorporate knowledge (inform it)
- Large Search Trees
 - How to make the tree smaller?
 - How to choose a search strategy?
 - How to choose an action to perform?
- Data structures + other optimizations


Heuristic Search

- Apply 'heuristic' guesses of value of nodes
- Node selection strategy
 - Move ordering
 - Iterative deepening
- "Pruning" nodes
- Bounding solutions



Heuristic Search Example

- Chess has ~40 moves at a state
- Most of them are bad
- Kasparov claimed to have only considered 2 moves per state
- Naïve search considers all 40
- Depth of search:
 - 40: 40, 1600, 64000, 2560000
 - 2: 2, 4, 8, 16, 32, depth 21
 - Pruning = Exponential gains



Recap

- **Problems** are defined by
 - Initial + Goal States, Actions, Cost Function
- **Optimal solution**
 - Has lowest cost function among all solutions
- **Search**
 - Finds solution by **exploring the search tree**
 - Search tree consists of **Nodes, Edges**
 - Nodes in search tree \neq states in problem
- Problem solving **performance** measured by
 - Completeness, Optimality, Space + Time Complexity